

DS-GA 3001.005
NYU Center for Data Science

Reinforcement Learning

Homework 02

Exercise 1 (20 points)

Markov Decision Processes and Bellman Equations:

- **1.1 (10 points):** What is a Markov Decision Process, and what is the Markov property?
Solution: An MDP is a mathematical idealization of goal-directed learning from interaction with an environment. Simulating an MDP produces a sequence of n tuples (trajectory) $(s_t, a_t, r_{t+1}, s_{t+1})_{t=0}^n = (s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_n)$. The MDP dynamics is fully characterized by the joint probability of each possible s_{t+1} and r_{t+1} as a function of the immediately preceding state and action, s_t and $a_t \Rightarrow p(s', r \mid s, a) = p(s_{t+1} = s', r_{t+1} = r \mid s_t = s, a_t = a)$.
Markov property: The state includes all information from past agent–environment interactions that influence the future $\Rightarrow p(s, r \mid s_t, a_t) = p(s, r \mid H_t, a_t)$.
- **1.2 (5 points):** Write $v_\pi(s)$ as a function of $q_\pi(s, a)$
Solution: $v^\pi(s) = \sum_a \pi(a \mid s)$ and $q^\pi(s, a) = E(q^\pi(s, a))$
- **1.3 (5 points):** Write $v_*(s)$ as a function of $q_*(s, a)$
Solution: $v^*(s) = \max_\pi v^\pi(s) = \max_a q^*(s, a)$

Exercise 2 (30 points)

Dynamic Programming:

- **2.1 (20 points):** In the gridworld problem below, the goal is to reach state g , the reward is -1 for moving to any state except state g where it is 0, actions in each state are up, down, right or left (by 1 step), and actions taking the agent off the grid leaves the state unchanged. What are the final state values after convergence of the *Value Iteration* algorithm?

		g	

Solution :

-4	-3	-2	-3
-3	-2	-1	-2
-2	-1	0	-1
-3	-2	-1	-2

We also accepted answers that took state value of 0 for states adjacent to G and so on

- **2.2 (5 points):** What is the key difference between the Policy Iteration algorithm and the Value Iteration algorithm?

Solution: VI is the same as PI except it truncates policy evaluation after just one loop (one update of each state). It is equivalent to turning the Bellman optimality equation into an update function

- **2.3 (5 points):** What is the key difference between synchronous Value Iteration and asynchronous Value Iteration?

Solution: Asynchronous DP backs up states in any order whatsoever, while synchronous DP proceeds with an entire loop over all possible states at each iteration

Exercise 3 (25 points)

Monte Carlo and Temporal Difference:

- **3.1 (5 points):** What is the key difference between Dynamic Programming and sample-based Reinforcement Learning?

Solution: We can accept any valid answer. Below are a few options:

Dynamic Programming requires a perfect model of state transitions and rewards to carry out a one-step look-ahead full-width backup at each iteration.

Sample-based RL uses a sample of experience to learn without a model. For example, instead of learning the true expected return $v^\pi(s) = E^\pi(G_t | s)$, we can sample its average:

$$v_{t+1}(s_t) = v_t(s_t) + \alpha_t \left(\sum_{k=0}^T \gamma^k r_{t+k+1} - v_t(s_t) \right)$$

- **3.2 (5 points):** List at least three differences between Monte-Carlo policy evaluation and Temporal Difference policy evaluation

Solution: We can accept any valid answer. Below are a few options:

- MC must wait until the end of the episode before the return is known.
- MC can only learn from complete sequences.
- MC can only work for episodic (terminating) environments.
- MC converges to the return $G_t = (r_{t+1} + \gamma(r_{t+2} + \dots))$ as an unbiased estimate of $v^\pi(s_t)$, but has high variance.
- MC converges to the best mean-squared fit for observed returns.
- MC does not bootstrap.
- TD can learn after every step, before knowing the final outcome.
- TD can learn from incomplete sequences.
- TD can work in continuing (non-terminating) environments.
- TD converges to the TD target $r_{t+1} + \gamma v(s_{t+1})$, which is a biased estimate of $v^\pi(s_t)$, but has lower variance.
- TD converges to the solution of the maximum likelihood Markov model.
- TD bootstraps.
- **3.3 (10 points):** Why is Q-learning considered an off-policy control method? Provide an example where using an off-policy method is more appropriate than using an on-policy method.
Solution: Because off-policy control means optimizing behavior by learning about a target policy from experience sampled by following a different policy called behaviour policy, and Q-learning systematically updates the greedy policy whatever the behavior policy followed is. If you want to explore a large action space efficiently, off-policy methods are better suited, for ex : Autonomous Driving
- **3.4 (5 points):** Why is Q-learning overly optimistic? And how does Double Q-learning help address this issue?
Solution: Q-learning uses the same Q-function to select greedy actions and to evaluate their values. This tends to over-select overestimated values and under-select underestimated values, perpetuating an upward bias. Double Q-learning uses two independent Q-functions to select vs. evaluate actions. It tends to reduce the upward bias of Q-learning estimation errors.

Exercice 4 (25 points)

Produce the code to do the following::

- **4.1 (10 points):** Write a function that implements the ϵ -greedy action-selection strategy. It should return the action selected. The required input to this function should be listed as arguments and/or defined in the function's docstring.

Solution: Two different examples of implementations are given inside lab 4

- **4.2 (15 points):** Write a function that implements a q-value update for tabular Q-learning, given a reward r observed after sampling an action a in a state s . It should return the updated q-value. The required input to this function should be listed as arguments and/or defined in the function's docstring.

Solution: Two different examples of implementations are given inside lab 4